

1. Explain what problem you are going to solve using this dataset. Provide a brief overview of your problem statement.

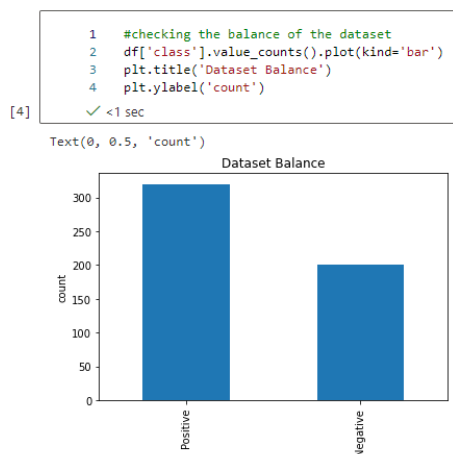
When diabetes is detected early, people can prevent progression and is less costly. If Diabetes goes undiagnosed or diagnosed too late, it can damage your heart, blood vessels , eyes, kidneys, and nerves. In 2018, 34.2 million Americans had diabetes. Of that 34.2 million people, 7.3 million were undiagnosed (<https://www.diabetes.org/>). To prevent undiagnosed Diabetes or late detection I am going to create an Early-Stage Diabetes classification model using this dataset.

The dataset will be checked for missing values (depending on how many of that specific feature is missing different approaches can be made – taking the average, removing those records and/or using the median value). After cleaning is done then preprocessing will take place. Here I will process the dataset to be a numerical and standardize the dataset if required based on the numerical values. Once that is completed then I will do feature selection to select the best features for the model. Using those selected features, I will train and test the model and use Accuracy as the success metric. The goal is to accurately predict whether someone has diabetes based on their answers to a survey of their symptoms/characteristics.

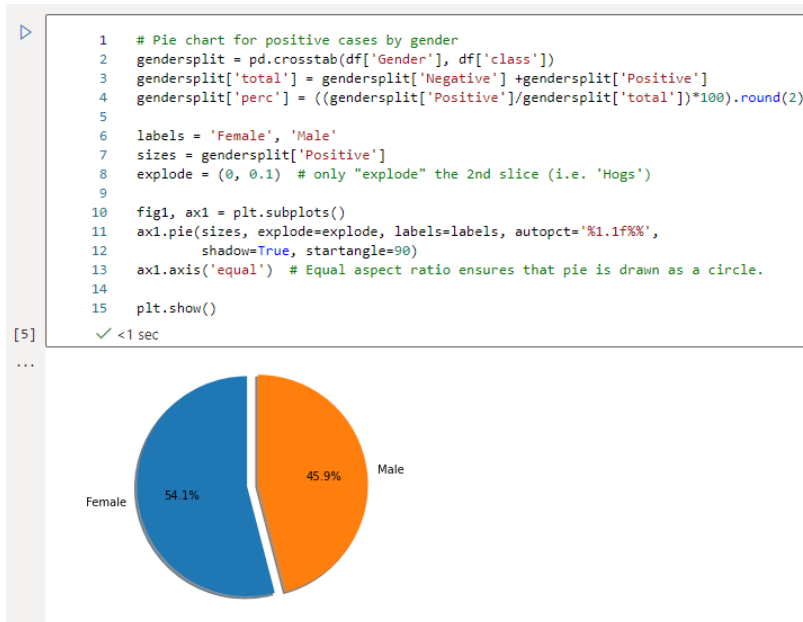
2. Explain your dataset. Explore your dataset and provide at least 5 meaningful charts/graphs with explanation.

This dataset consists of survey data of people between the ages 16-90, has 17 features (columns) and 520 records (rows). Only one column out of the whole dataset is not binary (Age column). The dataset consists of symptoms/characteristics that surgery respondents either responded with a ‘yes’ or a ‘no’ along with their age, gender and whether they do have diabetes or not.


The first graph that was created when exploring the dataset was a bar graph to show the balance of the dataset. By just looking at this graph one can notice that the dataset is unbalanced, there is a significant amount of more positive labelled cases than negative.



The second chart is a pie chart showing the number of people that have diabetes split by gender. When looking at the chart one can notice that there are more female positive cases than men. When extracting information like this from a chart its important to note the percentage of the number of males and females in the study to begin with. The total number of women in the study is 192 and 173 of them being positive meaning 90% of the women in the study have diabetes whereas for the men 147 out of 328 were positive for diabetes (45%).



The distribution graph and boxplot use the same data but are presented in different charts/graphs to show the different ways the same data can be presented. These graphs show the distribution of the survey takers age. You can see from the graph that the median age is just under 50 (the median in the boxplot is the line in the middle inside the box). From the boxplot you can also see the minimum, maximum, first quartile, third quartile and outliers.



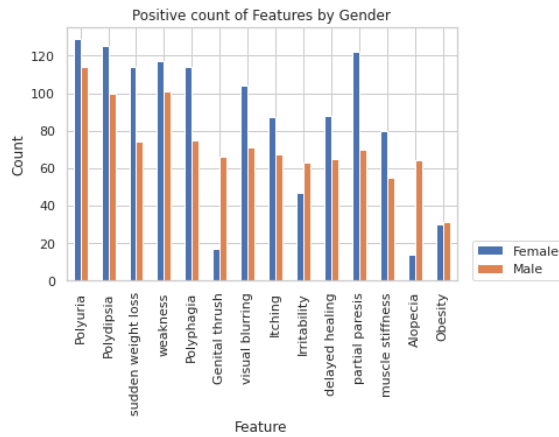
The bar graph below shows the comparison between the positive counts for diabetes between men and women by the features. This is a good way to see which features are more prominent in positive cases of women or men. For example, from the chart you can see that genital thrush and alopecia is much more common for diabetes in men in this dataset while partial paresis and sudden weight loss is more common with women.

```

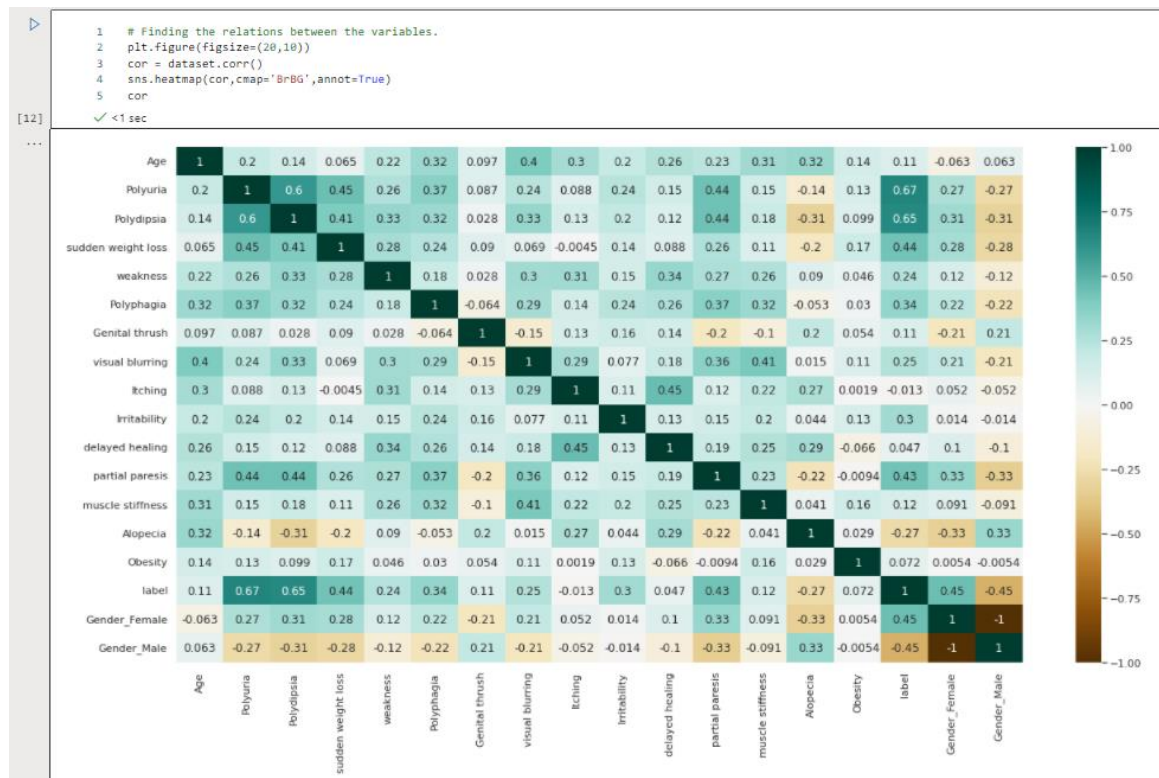
18
19 gender_features.plot(kind = 'bar')
20 plt.title('Positive count of Features by Gender')
21 plt.ylabel('Count')
22 plt.legend(loc=(1.04,0))
23 plt.xlabel('Feature')
24 plt.subplots_adjust(hspace=0.5)

```

[10] ✓ <1 sec



The final chart is a correlation heatmap. The heatmap visualizes the correlation between the different variables. From the heatmap one can see that the two main features affecting the label (target variable) are Polyuria and Polydipsia. The relationship between all the other variables can also be observed from the correlation heatmap.



3. Do data cleaning/pre-processing as required and explain what you have done for your dataset and why?

For data cleaning/pre-processing a few steps were taken. I started with checking for null values then encoding the dataset so that all string values (categorical values) would be numerical values. From there I did feature selection to select the best features for my model.

I first checked for null values to make sure that all records were complete. I found that there were no Null values so no further action was required to replace any values. Most of the columns as mentioned in the previous question were binary columns with strings “yes/no.” I converted all these columns to have 1/0 instead and used One Hot Encoding with the binary strategy for the gender column. The binary strategy essentially means that each category value was converted into a new column and assigned 1 or 0 based on that column encoded value. At this point all the columns are now numerical with all except the age column being binary columns. In this case, a lot of the data is binary, in these types of dataframes typically there is no need for scaling the data since most of the values are either 1 or 0.

Feature selection is used to minimize the number of input variables one needs for a model to predict the target variable. There are many different techniques for feature selection. The feature selection technique that was used in this model is Mutual Information with select K Best. Mutual information measures the dependencies between the variables. When the variables are independent from each other than the score is 0 and higher values mean higher dependency between the variables. The Select K best feature selection will then select the 10 best features based on the mutual information score.

```
[3] 1 #checking for missing values
    2 df.isnull().sum()
✓ <1 sec

... Age          0
    Gender       0
    Polyuria     0
    Polydipsia    0
    sudden weight loss 0
    weakness     0
    Polyphagia   0
    Genital thrush 0
    visual blurring 0
    Itching      0
    Irritability 0
    delayed healing 0
    partial paresis 0
    muscle stiffness 0
    Alopecia     0
    Obesity      0
    class        0
    dtype: int64
```

```
[8] 1 #encoding data to be 1/0 binary
    2 encode_df = df.replace(["Yes", 'No'],[1,0])
    3 encode_df['class'] = encode_df['class'].apply(lambda x: 1 if x == 'Positive' else 0)
✓ <1 sec
```

```
[9] 1 #oneHotEncoding gender column
    2 catecols = ['Gender']
    3 for col in catecols:
    4     dfeat = pd.get_dummies(encode_df[col], prefix=col, drop_first=False)
    5
    6     #adding the new feature to the dataframe
    7     encode_df = pd.concat([encode_df, dfeat], axis=1)
✓ <1 sec
```

```
[16] 1 #Feature selection techniques
    2 from sklearn.feature_selection import SelectKBest
    3 from sklearn.feature_selection import mutual_info_classif
    4
    5 # define Feature selection
    6 fs = SelectKBest(score_func = mutual_info_classif,k=10)
    7 #apply feature selection
    8 x = fs.fit_transform(x, y.ravel())
    9 print(x.shape)
✓ 1 sec

... (520, 10)
```

4. Implement 2 machine learning models, explain which algorithms you have selected and why. Compare them and show success metrics (Accuracy/RMSE/Confusion Matrix) as per your problem. Explain results.

The two machine learning algorithms that were chosen for this classification model is Support Vector Machine (SVM) and Decision Tree Classifier. SVM was chosen since this algorithm is known for performing well on datasets with binary classification and outliers having minimal impact. Decision Tree Classifier was also chosen due to popularity with binary classification and its ability to quickly classify unknown records. As patient records are very different from each other (every person has different health records) this feature was very important when choosing a model.

Comparison of the two models were done using accuracy and confusion matrix which can be seen in below screenshots. The confusion matrix shows True Positives (predicted positive and its true), True Negative (predicted negative and its true), False Positive (predicted positive and its false), and False Negative (predicted negative and its false). SVM had an accuracy of 0.929 and had TP of 47, FP of 7, FN of 4 and TN of 98. Decision Tree had an accuracy of 0.942 and had TP of 51, FP of 3, FN of 6 and TN of 96. Therefore, the better model between these two is the Decision Tree Classifier.

```
[24] def SVC(X_train, X_test, y_train, y_test):
      1
      2
      3 # Classification algorithm, predictions, and Success metrics.
      4 model = LinearSVC()
      5 model.fit(X_train, y_train)
      6 pred = model.predict(X_test)
      7 acc = accuracy_score(y_test, pred)
      8 cm = confusion_matrix(y_test, pred)
      9
     10 print('\nAccuracy of Support Vector Machine:', acc)
     11 print('\nConfusion Matrix of Support Vector Machine:\n', cm)
     12
     13 return model, pred, acc, cm
      ✓ <1 sec
```

```
[25] SVC = SVC(X_train, X_test, y_train, y_test)
      ✓ <1 sec
```

Accuracy of Support Vector Machine: 0.9294871794871795

Confusion Matrix of Support Vector Machine:

```
[[47  7]
 [ 4 98]]
```

```
[26] def DT(X_train, X_test, y_train, y_test):
      1
      2
      3 # Classification algorithm, predictions, and Success metrics.
      4 model = DecisionTreeClassifier(random_state=1)
      5 model.fit(X_train, y_train)
      6 pred = model.predict(X_test)
      7 acc = accuracy_score(y_test, pred)
      8 cm = confusion_matrix(y_test, pred)
      9
     10 print('\nAccuracy of Decision Tree:', acc)
     11 print('\nConfusion Matrix of Decision Tree:\n', cm)
     12
     13 return model, pred, acc, cm
      ✓ <1 sec
```

```
[27] DecTree = DT(X_train, X_test, y_train, y_test)
      ✓ <1 sec
```

Accuracy of Decision Tree: 0.9423076923076923

Confusion Matrix of Decision Tree:

```
[[51  3]
 [ 6 96]]
```

5. Use Automated ML for your data set. Explain best model results.

Below you can find the screenshots for the process of created an automated ML for my dataset. The process consisted of uploading my dataset, creating a cluster instance, selecting a classification algorithm, splitting the data and then running the model. In the screenshots you can see the different models that were run as well as their accuracy. The best model was found to be 'VotingEnsemble' which had an accuracy of 0.967. The top 4 features for this model were Polyuria, Polydipsia, Gender and Itching. Other success metrics can be noted in the screenshots below as well.

The screenshot shows the 'Create dataset from local files' wizard in Microsoft Azure Machine Learning Studio. The interface is divided into three main sections: a left-hand navigation pane, a central progress pane, and a right-hand configuration pane.

- Left-hand navigation pane:** Contains a list of tools and services including New, Home, Author, Notebooks, Automated ML (selected), Designer, Assets, Datasets, Experiments, Pipelines, Models, Endpoints, Manage, Compute, Environments, Datastores, Data Labeling, and Linked Services.
- Central progress pane:** Shows the steps of the wizard: Select dataset (active), Configure run, Select task and settings, [Optional] Validate and test, Basic info, Datastore and file selection (active), Settings and preview, Schema, and Confirm details.
- Right-hand configuration pane:**
 - Datastore and file selection:** A dropdown menu shows 'workspaceblobstore' as the selected datastore. Below it is a link to 'Create new datastore'.
 - Select files for your dataset:** A message states: 'These files will be uploaded to your selected datastore and made available in your workspace. Supported file types include: delimited (i.e. csv, tsv), Parquet, JSON Lines, and plain text.' Below this is an 'Upload' button and a status message: '1 files selected. Total size 0.03308 MiB. 0/1 files uploaded'.
 - File list:** A table displays the selected file:

File name	Size (MiB)	Upload %	Status
diabetes_data_upload.csv	0.03308		
 - Upload path:** A text field shows 'UI'. A tooltip message says: 'Files will be uploaded to: "\${Upload path}/11-30-2021_054742_UTC"'. Below the field is a checkbox for 'Skip data validation'.
- Bottom navigation:** Includes 'Back', 'Next', and 'Cancel' buttons.

The screenshot shows the 'Automated ML' page in Microsoft Azure Machine Learning Studio. It includes a header with navigation links and a main content area with a table of recent automated ML runs.

Let Automated ML train and find the best model based on your data without writing a single line of code. [Learn more about Automated ML](#)

+ New Automated ML run Refresh

Recent Automated ML runs								View all experiments
Display name	Experiment	Status	Submitted time	Duration	Submitted by	Compute target	Tags	
mighty_knee_z0vdys4d	diabetesprediction	Running	Nov 28, 2021 11:44 AM	15m 51s	Melina Fartaj	assign5		

Microsoft Azure Machine Learning Studio

Home > Automated ML > Start run

New

Home

Notebooks

Automated ML

Designer

Assets

Datasets

Experiments

Pipelines

Models

Endpoints

Manage

Compute

Environments

Datastores

Data Labeling

Linked Services

Create a new Automated ML run

Select dataset

Configure run

Select task and settings

[Optional] Validate and test

Configure run

Select from existing experiments or create a new experiment, then select the target column and training compute. [Learn more on how to configure the experiment.](#)

Dataset

Diabetes [\(View dataset\)](#)

Experiment name *

Create new

New experiment name

diabetesprediction

Target column *

class (String)

Select compute type

Compute instance

Select Azure ML compute instance *

assign5 - Running

+ New Refresh computes

Back

Next

Cancel

Microsoft Azure Machine Learning Studio

Home > Automated ML > Start run

New

Home

Notebooks

Automated ML

Designer

Assets

Datasets

Experiments

Pipelines

Models

Endpoints

Manage

Compute

Environments

Datastores

Data Labeling

Linked Services

Create a new Automated ML run

Select dataset

Configure run

Select task and settings

[Optional] Validate and test

Select task type

Select the machine learning task type for the experiment. To fine tune the experiment, choose additional configuration or featurization settings.

Classification

To predict one of several categories in the target column: yes/no, blue, red, green.

☐ Enable deep learning

Regression

To predict continuous numeric values

Time series forecasting

To predict values based on time

[View additional configuration settings](#)

[View featurization settings](#)

Back

Next

Cancel

Microsoft Azure Machine Learning Studio

HomeAutomated MLStart run

Create a new Automated ML run

Select dataset

Configure run

Select task and settings

Optional Validate and test

Optional Select the validation and test type

You can choose a validation type and select a test dataset as an optional step. Providing your own validation and test datasets are currently preview features.

Validation typeAuto

Test dataset (preview)Test split (choose a percentage of the training data)

Percentage test of data *30

Automated ML recommends that between 10 and 30 percent of data is held out for test

Back

Finish

Cancel

Learning Studio

HomeExperimentsdiabetespredictionmighty_knee_x0vdys4dmango_kitten_y61ffw0w

mango_kitten_y61ffw0w

RefreshDeployDownloadExplain modelTest model (preview)CancelDelete

DetailsModelExplanations (preview)MetricsData transformation (preview)Test results (preview)Outputs + logsImagesChild runsSnapshotMonitoring (preview)

Properties

StatusCompleted

CreatedNov 28, 2021 12:21 PM

StartedNov 28, 2021 12:21 PM

Duration1m 10.19s

Compute duration1m 10.19s

Compute targetassign5

Run IDAutoML_01261456-3ba5-4745-8048-b63dba24d75b_40

Script nameautoml_driver.py

Created byMelina Fartaj

Input datasetsInput name: training_data. Dataset: DiabetesDataSet: Version 1

Output datasetsNone

EnvironmentAzureML-AutoML90

ArgumentsNone

See all properties

Raw JSON

Tags

mflow.source.name : automl_driver.py

mflow.source.type : JOB

model_explain_run_id : AutoML_01261456-3ba5-4745-8048-b63dba24d75b_ModelExplain

model_explanation : True

Metrics

Accuracy0.96719

AUC macro0.99232

AUC micro0.99247

AUC weighted0.99232

Average precision score macro0.99111

Average precision score micro0.99277

View all other metrics

Description

Click edit icon to add a description

Learning Studio

Home

>

Automated ML

>

diabetesprediction

>

mighty_knee_z0vdy54d

mighty_knee_z0vdy54d

Refresh

Cancel

Delete

Details

Data guardrails

Models

Outputs + logs

Child runs

Snapshot

Refresh

Deploy

Download

Explain model

Edit columns

Reset view

Search

Submitted time

All filters

Clear all

Showing 1-25 of 42 models

Page size: 25

Algorithm name	Explained	Accuracy ↓	Sampling	Submitted time	Duration	Hyperparameter
VotingEnsemble	View explanation	0.96719	100.00 %	Nov 28, 2021 12:21 PM	1m 10s	algorithm : ['SVM', 'LightGBM', 'XG']
StackEnsemble		0.96164	100.00 %	Nov 28, 2021 12:22 PM	1m 14s	algorithm : ['SVM', 'LightGBM', 'XG']
StandardScalerWrapper, SVM		0.95623	100.00 %	Nov 28, 2021 11:48 AM	22s	C : 16.768329368110066 class :
MaxAbsScaler, LightGBM		0.95616	100.00 %	Nov 28, 2021 11:47 AM	33s	min_data_in_leaf : 20
StandardScalerWrapper, XGBoostClassifier		0.95601	100.00 %	Nov 28, 2021 12:09 PM	47s	booster : gbtree colsample_by:
MaxAbsScaler, XGBoostClassifier		0.95345	100.00 %	Nov 28, 2021 11:47 AM	22s	tree_method : auto
SparseNormalizer, XGBoostClassifier		0.95330	100.00 %	Nov 28, 2021 12:14 PM	51s	booster : gbtree colsample_by:
RobustScaler, KNN		0.95083	100.00 %	Nov 28, 2021 11:48 AM	22s	metric : manhattan n_neighbors :
SparseNormalizer, XGBoostClassifier		0.95068	100.00 %	Nov 28, 2021 12:11 PM	52s	booster : gbtree colsample_by:
MinMaxScaler, KNN		0.94812	100.00 %	Nov 28, 2021 11:48 AM	22s	metric : manhattan n_neighbors :

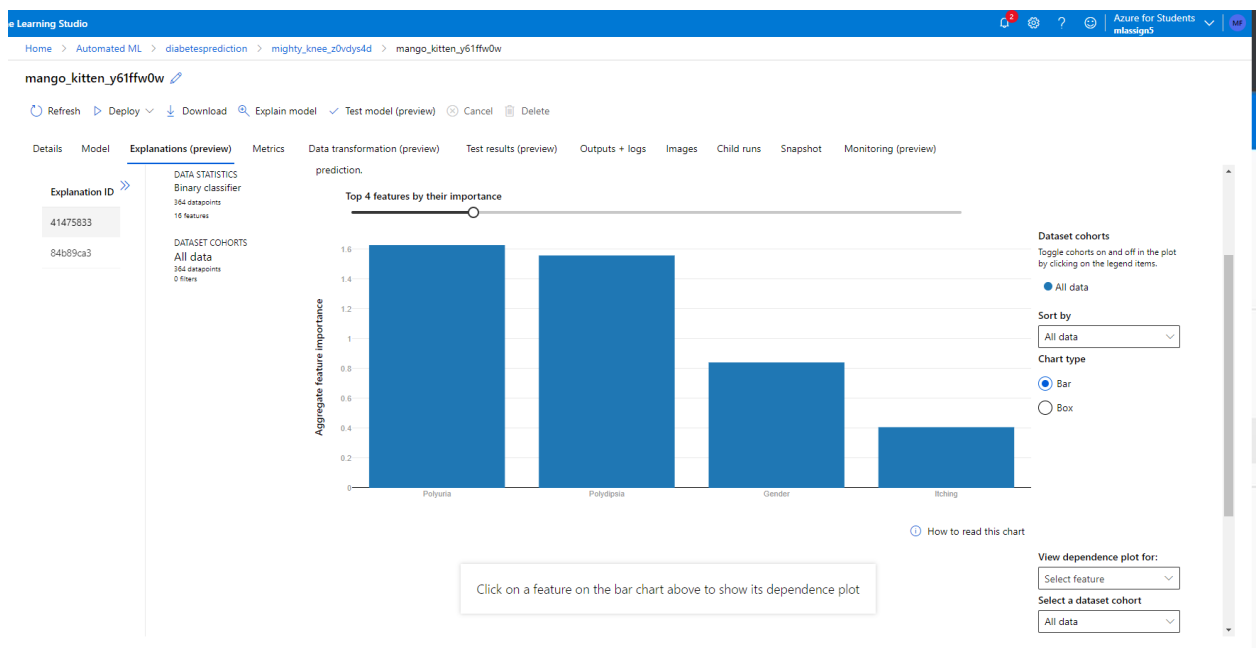
<<

<

Page 1 of 2

>

>>



Learning Studio

Home > Automated ML > diabetesprediction > mighty_knee_r0vdy54d > mango_kitten_y61ffw0w

mango_kitten_y61ffw0w

Refresh Deploy Download Explain model Test model (preview) Cancel Delete

Details Model Explanations (preview) **Metrics** Data transformation (preview) Test results (preview) Outputs + logs Images Child runs Snapshot Monitoring (preview)

Select a metric to see a visualization or table of the data.

Search

- ☒ accuracy
- ☒ accuracy_table
- ☒ AUC_macro
- ☒ AUC_micro
- ☒ AUC_weighted
- ☐ average_precision_score_macro
- ☐ average_precision_score_micro
- ☐ average_precision_score_weighted
- ☐ balanced_accuracy
- ☐ confusion_matrix
- ☐ f1_score_macro
- ☐ f1_score_micro
- ☐ f1_score_weighted
- ☐ log_loss
- ☐ matthews_correlation

View as: ☒ Chart ☐ Table

accuracy	AUC_macro	AUC_micro	AUC_weighted
0.967	0.992	0.992	0.992

Precision-Recall

ROC

Azure for Students mlassign5

Ensemble details

Select an ensemble algorithm to see the ensemble weights and hyperparameters.

- ☒ SparseNormalizer, XGBoostClas...
- ☐ StandardScalerWrapper, XGBoo...
- ☐ RobustScaler, KNN
- ☐ StandardScalerWrapper, SVM
- ☐ MaxAbsScaler, XGBoostClassifier
- ☐ MaxAbsScaler, LightGBM

Ensemble weight: 0.16666666666666666

Data transformation:

```
1 {
2   "class_name":
3   "SparseNormalizer",
4   "module": "automl.
5   client.core.common.
6   model_wrappers",
7   "param_args": [],
8   "param_kwargs": {
9     "norm": "l2"
10 }
```

Training algorithm:

```
1 {
2   "class_name":
3   "XGBoostClassifier",
4   "module": "automl.
5   client.core.common.
6   model_wrappers",
7   "param_args": [],
8   "param_kwargs": {
9     "booster": "gbtree",
10    "colsample_bytree":
11    0.7,
12    "eta": 0.1,
13    "max_depth": 4,
14    "max_leaves": 0,
15    "n_estimators": 100,
16    "objective":
17    "reg:logistic",
18    "reg_alpha": 0.
19 }
```